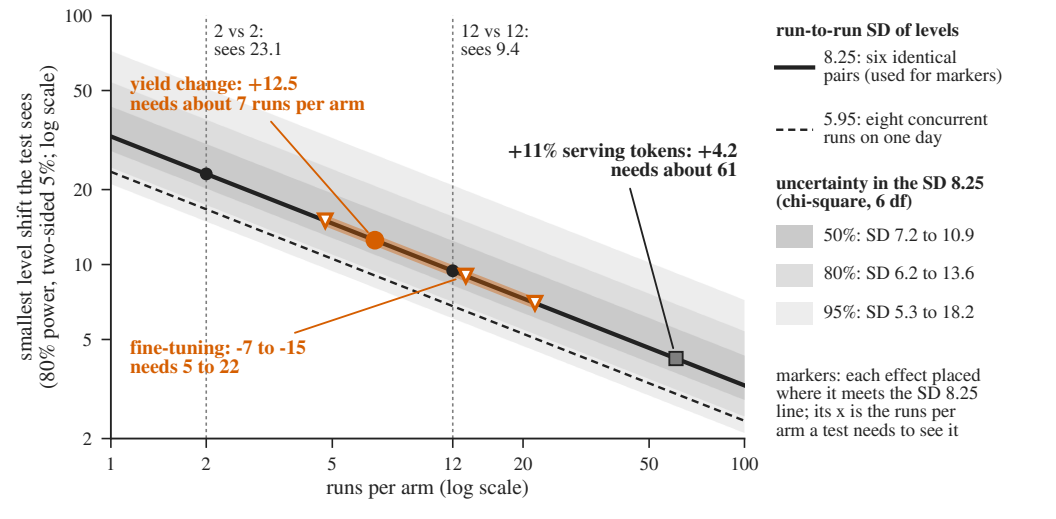
smallest level shift the test sees
(80% power, two-sided 5%; log scale)
100
50
20
10
5
2
1
2
5
12
20
50
100
runs per arm (log scale)
2 vs 2:
sees 23.1
12 vs 12:
sees 9.4
yield change: +12.5
needs about 7 runs per arm
+11% serving tokens: +4.2
needs about 61
fine-tuning: -7 to -15
needs 5 to 22
run-to-run SD of levels
8.25: six identical pairs (used for markers)
5.95: eight concurrent runs on one day
uncertainty in the SD 8.25 (chi-square, 6 df)
50%: SD 7.2 to 10.9
80%: SD 6.2 to 13.6
95%: SD 5.3 to 18.2
markers: each effect placed where it meets the SD 8.25 line; its x is the runs per arm a test needs to see it